

## What is a Terraform Provider?

A **Terraform Provider** is a plugin that enables Terraform to manage and interact with external APIs and services (like AWS, Azure, GitHub, Kubernetes, etc.). Providers expose resources and data sources, which are used in Terraform configurations to describe infrastructure or platform elements declaratively.

---

### Core Concepts of Terraform Providers

#### 1. Plugin Architecture

- Providers are distributed as Go binaries (plugins).
- When terraform init runs, it downloads required providers into the .terraform directory.
- Providers communicate with APIs over HTTP, gRPC, or CLI interfaces (e.g., kubectl, az, aws).

#### 2. Authentication and Configuration

- Most providers support authentication via environment variables, credentials files, or explicit block configuration.
- Configuration block often includes:

```
provider "aws" {  
    region = "us-west-2"  
    profile = "devops"  
}
```

#### 3. Versioning

- Providers follow semantic versioning (x.y.z) and are declared like:

```
terraform {  
    required_providers {  
        aws = {  
            source = "hashicorp/aws"  
            version = "~> 5.0"  
        }  
    }  
}
```

```
}
```

```
}
```

---

## Commonly Used Terraform Providers (with Use Cases)

Provider	Source	Use Case
aws	hashicorp/aws	Provision EC2, VPCs, S3, IAM, EKS
azurerm	hashicorp/azurerm	Manage Azure resources like VMs, ACR, AKS
google	hashicorp/google	GCP: GKE, Compute Engine, IAM, Cloud SQL
kubernetes	hashicorp/kubernetes	Apply Kubernetes manifests declaratively
helm	hashicorp/helm	Manage Helm charts within Kubernetes
vault	hashicorp/vault	Secrets, policies, auth methods
github	integrations/github	Manage repos, teams, webhooks, permissions
datadog	DataDog/datadog	Infrastructure and application monitoring
gitlab	gitlabhq/gitlab	GitLab projects, groups, CI/CD variables
newrelic	newrelic/newrelic	Application monitoring and alerting
cloudinit	hashicorp/cloudinit	Cloud-init configuration for compute instances
docker	kreuzwerker/docker	Manage Docker containers, images, and networks
random	hashicorp/random	Generate random values for secrets or names
tls	hashicorp/tls	Generate TLS certificates and keys
null	hashicorp/null	Run provisioners and custom scripts
http	hashicorp/http	Call APIs or perform health checks

---

## Advanced Usage Patterns

### 1. Multiple Providers / Aliases

Use aliases to configure the same provider for different environments:

```
provider "aws" {  
    alias = "prod"  
  
    region = "us-east-1"  
}  
  
provider "aws" {  
    alias = "dev"  
  
    region = "us-west-2"  
}
```

Then reference using:

```
resource "aws_s3_bucket" "prod_bucket" {  
  
    provider = aws.prod  
  
    bucket  = "prod-bucket"  
}
```

---

### 2. Chained or Nested Providers

- For layered platforms (e.g., Terraform Helm provider + Kubernetes provider + AWS EKS cluster).
- Example:

```
provider "kubernetes" {  
  
    host          = aws_eks_cluster.mycluster.endpoint  
  
    cluster_ca_certificate = base64decode(aws_eks_cluster.mycluster.certificate_authority.0.data)  
  
    token         = data.aws_eks_cluster_auth.mycluster.token  
}
```

```
resource "helm_release" "nginx" {  
    name      = "nginx"  
    repository = "https://charts.bitnami.com/bitnami"  
    chart      = "nginx"  
}
```

---

## Creating a Custom Provider (For Plugin Developers)

- Providers are written in Go using the [Terraform Plugin SDK](#).
- Structure includes:
  - Provider() function
  - schema.Provider definition
  - ResourcesMap and DataSourcesMap
- Example repo: [terraform-provider-example](#)

---

## Best Practices for Using Providers

1. **Lock versions** to ensure consistent behavior across environments.
2. **Avoid sensitive data** in provider blocks; use environment variables or secrets managers.
3. **Use terraform validate and terraform plan** to catch errors early.
4. **Combine providers** logically (e.g., AWS + Kubernetes + Helm for EKS-based app deployments).
5. **Modularize your code** and pass providers as inputs using providers = { aws = aws.dev }.

---

## Example: Secure Setup for Vault and AWS

```
provider "vault" {  
    address = "https://vault.company.com"
```

```
token = var.vault_token  
}
```

```
provider "aws" {  
  region = "us-west-2"  
  profile = "devops"  
}
```

Use Vault to dynamically generate AWS credentials instead of static ones.

## Terraform Resources

A **resource** in Terraform represents a piece of infrastructure you want to create, manage, or destroy. It's the "**thing you own or manage**" in your infrastructure—like an EC2 instance, a Kubernetes deployment, a VPC, a DNS record, etc.

### Syntax:

```
resource "<PROVIDER>_<TYPE>" "<NAME>" {  
  # configuration arguments  
}
```

---

### Example: AWS EC2 Instance

```
resource "aws_instance" "web_server" {  
  ami      = "ami-0c55b159cbfafe1f0"  
  instance_type = "t2.micro"  
  tags = {  
    Name = "WebServer"  
  }  
}
```

```
resource "aws_instance" "web_server1" {
    ami      = "ami-0c55b159cbfafe1f0"
    instance_type = "t2.micro"
    tags = {
        Name = "WebServer"
    }
}
```

This creates a t2.micro EC2 instance in AWS with the specified AMI.

---

### Example: Azure Resource Group

```
resource "azurerm_resource_group" "main" {
    name    = "devops-rg"
    location = "East US"
}
```

---

### Terraform Data Sources

**Data sources** allow Terraform to **query and read external information** (read-only) so you can use that information in your resources. They **do not create anything**, but they are useful for **reusing existing infrastructure**.

#### Use Case:

You're deploying a VM into a VPC that's already created. Instead of hardcoding the VPC ID, you query it dynamically using a data block.

---

### Example: Get Existing AWS VPC by Tag

```
data "aws_vpc" "selected" {
    filter {
```

```
    name  = "tag:Name"
    values = ["main-vpc"]
}
}
```

This retrieves details about an existing VPC with the tag Name=main-vpc.

You can then use it like:

```
resource "aws_subnet" "example" {
  vpc_id  = data.aws_vpc.selected.id
  cidr_block = "10.0.1.0/24"
}
```

---

#### Example: Get Latest Ubuntu AMI from AWS

```
data "aws_ami" "ubuntu" {
  most_recent = true
  owners      = ["099720109477"] # Canonical
  filter {
    name  = "name"
    values = ["ubuntu/images/hvm-ssd/ubuntu-focal-20.04-amd64-server-*"]
  }
}
```

Used with:

```
resource "aws_instance" "ubuntu_server" {
  ami        = data.aws_ami.ubuntu.id
  instance_type = "t2.micro"
}
```

## Comparison: resource vs data

Feature	Resource	data
Action	Creates or manages infrastructure	Reads existing infrastructure
Lifecycle	Can be created, updated, destroyed	Read-only
Example	EC2 instance, S3 bucket, VM	VPC lookup, AMI lookup, Key Vault secret
Useful for	Provisioning	Referencing external state

---

## Advanced Use Cases for DevOps Engineers

### 1. Use Data Source to Fetch AWS Secrets from Secrets Manager

```
data "aws_secretsmanager_secret_version" "db_password" {  
    secret_id = "prod/db-password"  
}  
  
resource "aws_db_instance" "mysql" {  
    # ...  
    password = data.aws_secretsmanager_secret_version.db_password.secret_string  
}
```

---

### 2. Use data for Modular Inputs

In a multi-account, shared VPC architecture:

```
data "terraform_remote_state" "network" {  
    backend = "s3"  
    config = {  
        bucket = "network-state"  
        key    = "vpc/terraform.tfstate"  
    }  
}
```

```
region = "us-west-2"

}

}

resource "aws_instance" "app" {

  subnet_id = data.terraform_remote_state.network.outputs.private_subnet_id
}
```

---

## ❖ Best Practices for Experienced DevOps Engineers

1. Use data sources to avoid hardcoding IDs or values.
  2. Combine data and resource for reusable, modular code.
  3. Avoid using data for resources you plan to manage with Terraform.
  4. Use `terraform_remote_state` for inter-module communication.
  5. Output values from modules and access them via data.
- 

## 🚀 Summary

Term	Purpose	Example Use Case
resource	Creates and manages infrastructure	Launch an EC2, create S3 bucket, etc.
data	Reads external or existing resources	Get existing VPC, latest AMI, secrets, etc.

## What is `terraform.lock.hcl`?

The `terraform.lock.hcl` file (also known as the *dependency lock file*) is a **generated file** that records the exact versions of Terraform providers used in your configuration. It ensures **repeatable, deterministic builds** of your infrastructure across teams, machines, and CI/CD environments.

---

## Why is `terraform.lock.hcl` Important?

Terraform configurations often define **version constraints** like:

```
terraform {  
  required_providers {  
    aws = {  
      source  = "hashicorp/aws"  
      version = "~> 5.0"  
    }  
  }  
}
```

This allows **flexibility**, but during initialization (`terraform init`), Terraform will download the **latest matching version**—which could change tomorrow. That's risky.

So, Terraform generates a `terraform.lock.hcl` file that **pins the exact provider version**, platform, and SHA checksum:

- Consistency across environments
  - Better security by verifying checksums
  - Eliminates version drift in teams and pipelines
- 

## Where Is `terraform.lock.hcl` Stored?

- Located in the **root directory** of your Terraform project (same place as `main.tf`, `variables.tf`).
  - Automatically created/updated by `terraform init`.
- 

## Sample `terraform.lock.hcl` Content

```
provider "registry.terraform.io/hashicorp/aws" {  
  version  = "5.13.0"
```

```
constraints = "~> 5.0"

hashes = [
    "h1:b18W6ly...",
    "zh:sha256:9f0a..."
]

}
```

#### Breakdown:

Field	Description
provider	Fully qualified name of the provider from Terraform Registry
version	The exact version downloaded and locked
constraints	Version constraint from your terraform block
hashes	SHA-256 digests used to validate provider binaries (for security)

---

#### How it Works in Practice

1.  You write:

```
terraform {

  required_providers {

    aws = {

      source  = "hashicorp/aws"
      version = "~> 5.0"
    }
  }
}
```

2.  Run terraform init:

- o Downloads latest version matching  $\sim> 5.0$  (e.g., 5.13.0)

- Generates `terraform.lock.hcl` with full hash and version
3. 📦 When others clone your repo and run `terraform init`:
- Terraform reads `terraform.lock.hcl`
  - Downloads **same exact version** of each provider
  - Verifies binary checksum with hashes

## ⌚ When Does It Change?

- If you change version constraints or add a new provider and run `terraform init`.
- You can force an update:

```
terraform init -upgrade
```

---

## 🔒 Security & Integrity

- Hash verification ensures **binary tampering detection**.
  - Terraform won't proceed if downloaded binary's hash doesn't match.
- 

## 🤝 Collaboration Best Practices

Practice	Why
Commit <code>terraform.lock.hcl</code> to VCS	Ensures everyone uses the same provider versions
Avoid manual edits	File is machine-generated and changes should be tracked via Terraform commands
Use <code>-upgrade</code> explicitly	Prevents accidental provider upgrades
Keep it versioned per workspace/environment	Different environments might use different providers/versions

---

## ⚠ Common Misunderstandings

Myth	Truth
You must write <code>terraform.lock.hcl</code>	✗ It's generated automatically
Lock file is optional	✓ Technically optional, but essential for team-based workflows
It locks modules	✗ It only locks <b>providers</b> , not Terraform modules (use Git commits/tags for that)

### Useful Commands

Command	Purpose
<code>terraform init</code>	Creates or updates <code>terraform.lock.hcl</code>
<code>terraform init -upgrade</code>	Updates to latest allowed provider versions
<code>terraform providers lock</code>	Manually regenerate lock file
<code>terraform providers mirror</code>	Download providers to local filesystem for air-gapped installs

### Summary

- `terraform.lock.hcl` is like a `package-lock.json` (Node.js) or `Pipfile.lock` (Python) but for **Terraform providers**.
- Ensures infrastructure deployments are **reproducible, secure, and collaborative**.
- Should always be **committed to version control** and never edited manually.

Upgrading the **AWS Terraform provider in a production environment** must be done **safely, deliberately, and in a controlled way** to prevent disruptions to live infrastructure.

Below is a **detailed, step-by-step guide for experienced DevOps engineers** on how to upgrade the AWS provider in production using Terraform.

### Goals of Provider Upgrade in Production

- Upgrade the AWS provider version (e.g., from 4.x to 5.x)
  - Avoid service interruptions
  - Maintain compatibility with existing Terraform code
  - Ensure idempotent and safe terraform apply
  - Work across multiple environments (dev, staging, prod)
- 

## Prerequisites

1. **Version control (Git)** with infrastructure code
  2. **CI/CD pipeline or Terraform workspace separation**
  3. **Terraform state backend configured** (e.g., S3 + DynamoDB)
  4. **Backup current state** before any upgrade
- 

## High-Level Steps

1. Review Provider Upgrade Guide and Changelog
  2. Update required\_providers version in `terraform {}` block
  3. Run `terraform init -upgrade` in **dev/staging**
  4. Perform **terraform plan** and fix any breaking changes
  5. Test in **non-production environments**
  6. Commit updated `terraform.lock.hcl` and code to version control
  7. Repeat plan + apply in **production**
- 

## Step-by-Step Instructions

### Step 1: Review AWS Provider Changelog

- Visit: <https://registry.terraform.io/providers/hashicorp/aws/latest/docs/guides/version-5-upgrade>
- Understand **breaking changes** (e.g., removal of deprecated attributes, behavior changes)

- Search for changes related to resources you're using: EC2, IAM, VPC, EKS, etc.
- 

## Step 2: 🖍️ Update Required Provider Block in main.tf

Update the version constraint:

```
terraform {  
    required_providers {  
        aws = {  
            source  = "hashicorp/aws"  
            version = "~> 5.0" # was "~> 4.0"  
        }  
    }  
}
```

**Tip:** Use  $\geq 5.0, < 6.0$  for more control.

---

## Step 3: 💡 Upgrade in Development Environment

```
terraform init -upgrade
```

- This downloads the latest compatible version of AWS provider
  - It will regenerate the `terraform.lock.hcl` file with updated version and hash
- 

## Step 4: 📋 Run Terraform Plan

```
terraform plan
```

- Carefully review the output
- Look for changes to resources that **should not be touched**
- Fix issues related to deprecated attributes or required new ones

### ✓ Best Practice:

- Use the `TF_LOG=DEBUG` env variable for deep debugging

- Run terraform validate for static analysis
- 

### Step 5: 🔎 Test Thoroughly in Staging

- Deploy using terraform apply in staging.
  - Test all critical infrastructure workflows (e.g., API response, autoscaling, data access).
  - Use monitoring tools (CloudWatch, Datadog, NewRelic) to confirm stability.
- 

### Step 6: ✅ Commit and Push Code

```
git add .
```

```
git commit -m "Upgrade AWS Terraform provider to v5.13.0"
```

```
git push origin main
```

Ensure `terraform.lock.hcl` is committed so other environments and CI/CD pipelines use the same version.

---

### Step 7: 🚨 Upgrade in Production Environment

Use the same Git branch and lock file:

```
bash
```

```
terraform init
```

```
terraform plan
```

```
terraform apply
```

🔴 **Do not** use `-upgrade` in production — rely on the committed lock file.

Use a **manual approval step** in CI/CD pipeline or human validation before applying.

---

### ✳️ Example GitHub Actions Snippet

```
yaml
```

```
jobs:
```

```
terraform-apply-prod:

  name: Terraform Apply (Prod)

  runs-on: ubuntu-latest

  environment: production

  steps:

    - uses: actions/checkout@v3

      - name: Setup Terraform

        uses: hashicorp/setup-terraform@v2

    - name: Terraform Init

      run: terraform init

    - name: Terraform Plan

      run: terraform plan -out=tfplan

    - name: Manual Approval

      uses: hmarr/auto-approve-action@v2

      with:

        approve: false # Manual input required

    - name: Terraform Apply

      run: terraform apply tfplan
```

---

 **Optional: Lock Specific Provider Versions for Environments**

Use different provider versions for dev, qa, and prod using workspaces or environment-specific terraform blocks.

---

## Rollback Strategy

Terraform doesn't natively support provider **downgrade rollback**. You should:

- Keep a backup of `terraform.lock.hcl` before upgrading
  - Revert code and lock file in Git
  - Run `terraform init` with the old lock file
  - Never apply if you suspect state corruption — validate first!
- 

## Summary: Provider Upgrade Checklist

Task	Status
Review AWS provider changelog	
Update required_providers version	
Run <code>terraform init -upgrade</code> in dev	
Fix deprecated/changed attributes	
Validate and test in staging	
Commit <code>terraform.lock.hcl</code>	
Apply in production with approval	

## What are Terraform Input Variables?

Terraform **input variables** allow you to **parameterize your code**—so it's **reusable, modular, environment-independent**, and **less error-prone**.

They're defined in your Terraform module and passed values via:

- \*.tfvars files
  - Environment variables
  - -var flag on CLI
  - terraform.tfvars
  - Variable blocks in main.tf
- 

## Syntax of a Variable

```
variable "instance_type" {  
    description = "EC2 instance type"  
    type        = string  
    default     = "t2.micro"  
}
```

**Usage:**

```
resource "aws_instance" "app" {  
    ami      = var.ami_id  
    instance_type = var.instance_type  
}
```

---

## Variable Types and Their Use in AWS Context

Type	Example Usage (AWS)
string	EC2 instance type, region, AMI ID
number	Number of EC2 instances, timeout seconds
bool	Enable/disable a feature like auto-scaling

Type	Example Usage (AWS)
list()	List of subnets, AZs, or security group IDs
map()	Tag values, AMI per region
object()	Complex input for multi-AZ or layered config
any	Accepts any type (use carefully)

---

## 🔥 AWS Use Case Examples for Input Variables

---

### ◆ 1. EC2 Launch with Parameterized Variables

**variables.tf**

```
variable "region" {
  type    = string
  description = "AWS region"
  default  = "us-east-1"
}
```

```
variable "ami_id" {
  description = "AMI ID to use"
  type    = string
}
```

```
variable "instance_type" {
  type    = string
  default = "t3.micro"
}
```

## main.tf

```
provider "aws" {  
    region = var.region  
}  
  
resource "aws_instance" "web" {  
    ami      = var.ami_id  
    instance_type = var.instance_type  
    tags = {  
        Name = "WebServer"  
    }  
}
```

## terraform.tfvars

```
ami_id = "ami-0c55b159cbfafe1f0"
```

---

### ◆ 2. Use List to Deploy into Multiple Subnets

```
variable "subnet_ids" {  
    description = "List of subnet IDs"  
    type       = list(string)  
}  
  
resource "aws_instance" "multi_az_app" {  
    count      = length(var.subnet_ids)  
    ami        = var.ami_id  
    instance_type = "t3.micro"  
    subnet_id   = var.subnet_ids[count.index]
```

```
}
```

---

### ◆ 3. Map for Region-wise AMI Lookup

```
variable "region_ami_map" {  
    description = "Map of AMI per region"  
    type      = map(string)  
    default = {  
        "us-east-1" = "ami-abc123"  
        "us-west-1" = "ami-def456"  
    }  
}  
  
provider "aws" {  
    region = var.region  
}  
  
resource "aws_instance" "web" {  
    ami      = var.region_ami_map[var.region]  
    instance_type = "t3.micro"  
}
```

---

### ◆ 4. Object for Advanced Resource Control

```
variable "ec2_config" {  
    description = "EC2 configuration"  
    type = object({  
        ami_id      = string  
        instance_type = string  
        enable_monitoring = bool
```

```
    })
}

resource "aws_instance" "custom" {
    ami          = var.ec2_config.ami_id
    instance_type = var.ec2_config.instance_type
    monitoring   = var.ec2_config.enable_monitoring
}
```

---

#### ◆ 5. Conditional Logic Using Boolean

```
variable "enable_ebs" {
    type  = bool
    default = false
}

resource "aws_ebs_volume" "extra" {
    count = var.enable_ebs ? 1 : 0
    availability_zone = "us-east-1a"
    size      = 10
}
```

---

#### ◆ 6. Use Variable with Default Tags Across Resources

```
variable "common_tags" {
    type = map(string)
    default = {
        Environment = "dev"
        Owner      = "DevOpsTeam"
    }
}
```

```

}

resource "aws_instance" "app" {

  ami      = var.ami_id

  instance_type = "t3.micro"

  tags      = var.common_tags

}

```

---

### Tips for Experienced DevOps Engineers

Tip	Description
 <b>Type Safety</b>	Always define the type of variables—especially in modules
 <b>Reusability</b>	Use variables to build reusable modules for VPCs, EC2, RDS, etc.
 <b>Secrets Management</b>	Never store sensitive data in variables. Use environment vars or vaults
 <b>Dynamic Inputs</b>	Use terraform.workspace and locals to differentiate per environment
 <b>Automate via CI/CD</b>	Pass variables using .tfvars or command line in Jenkins, GitHub Actions
 <b>Variable Validation</b>	Use validation block for input checks (Terraform >= 0.13)

**Example:**

```

variable "instance_type" {

  description = "EC2 instance type"

  type      = string

  validation {

    condition  = contains(["t3.micro", "t3.small", "t3.medium"], var.instance_type)

    error_message = "Invalid instance type. Use t3.micro, t3.small, or t3.medium"

  }
}

```

}

## ✓ Summary

Feature	Benefit
Input variables	Parameterize your Terraform for reusability
string	Regions, instance types, AMI IDs
list	Multi-subnet, multi-AZ setups
map	AMIs per region, global tag maps
object	Group related config like EC2 specs
validation	Enforce safe, allowed inputs

In Terraform, **variables** are a powerful way to parameterize configurations and separate code from environment-specific values. There are **multiple ways** to assign values to variables, and Terraform follows a strict **precedence order** to determine which values are used.

Let's break it all down clearly — including a **focus on `terraform.tfvars`**, which is commonly used in production-grade DevOps setups.

---

## ✓ Declaring Variables

First, you define variables in a `variables.tf` file (or in any `.tf` file):

```
variable "region" {  
    description = "AWS region"  
    type        = string  
    default     = "us-west-2" # Optional  
}
```

---

## ⌚ 7 Ways to Assign Values to Variables

Terraform assigns variable values from **multiple sources**, in the following **precedence order** (highest to lowest):

Priority	Method	Description
1	<b>CLI -var flag</b>	Highest priority — used for overrides and testing
2	<b>CLI -var-file flag</b>	Explicit file passed with terraform plan/apply
3	<b>Environment variables</b>	TF_VAR_name style
4	<b>terraform.tfvars file</b>	Automatically loaded if present
5	<b>*.auto.tfvars files</b>	Auto-loaded if they exist
6	<b>Default values</b> in variable blocks	Fallback values
7	<b>Module input arguments</b>	When calling modules with module.<name>.inputs

## Variable Assignment Methods — Detailed

### 1 CLI Flag (-var)

```
terraform plan -var="region=us-east-1"
```

- Great for **quick tests**
  - Avoid using in CI/CD or production (less traceable)
- 

### 2 CLI -var-file

```
terraform apply -var-file="prod.tfvars"
```

- You can maintain different .tfvars files per environment
- Supports HCL or JSON format

Example prod.tfvars:

```
region = "us-east-1"  
instance_type = "t3.medium"
```

---

### 3 Environment Variables (TF\_VAR\_)

```
export TF_VAR_region="ap-south-1"
```

terraform plan

- Used in CI/CD pipelines (e.g., GitHub Actions, Jenkins)
  - Good for injecting **secrets** or credentials (but be careful)
- 

### 4 terraform.tfvars file (Primary Focus)

- Auto-loaded **by default** when running any Terraform command
- Located in the **root** of your working directory
- No need to pass with -var-file

Example terraform.tfvars:

```
region    = "eu-west-1"  
instance_type = "t2.micro"
```

#### ✓ Best Practice:

- Use this for **default project-level values**
  - Keep this **environment-specific**, but generic enough to be version-controlled
- 

### 5 \*.auto.tfvars files

- Auto-loaded just like terraform.tfvars, but you can have **multiple files**
- Great for **layered configuration**: env.auto.tfvars, backend.auto.tfvars

Example:

```
production.auto.tfvars
```

```
database.auto.tfvars
```

⚠ If multiple \*.auto.tfvars are present, they are **merged in alphanumeric order**

---

## 6 Default Values in variable Block

```
variable "region" {  
    type = string  
    default = "us-west-2"  
}
```

- Used only if no other value is provided
  - Great for **sane defaults**, but don't rely on it in production
- 

## 7 Module Inputs

If you're calling a module:

```
module "vpc" {  
    source = "./modules/vpc"  
    region = var.region  
}
```

You pass values via `module.<name>` block; they can also be dynamically passed from another variable.

---

## 8 Precedence Order in Action

Let's say `region` is defined in all the following locations. Which one wins?

```
variable "region" {  
    default = "us-west-1"  
}  
  
CLI flag: -var="region=us-east-1" ✅ (Wins)  
  
CLI var file: prod.tfvars → region = "us-west-2"  
  
terraform.tfvars → region = "eu-central-1"
```

```
TF_VAR_region="ap-south-1"
```

```
Default = "us-west-1"
```

 Terraform chooses the **first match in precedence order**.

---

## Real-World Folder Structure

```
infra/
    ├── main.tf
    ├── variables.tf
    ├── terraform.tfvars      # Default project-wide vars
    ├── prod.tfvars          # Used with -var-file for prod
    ├── dev.auto.tfvars       # Auto-loaded only for dev
    └── outputs.tf
```

---

## Best Practices for Using terraform.tfvars

Practice	Benefit
<input checked="" type="checkbox"/> Keep it in Git	Makes configurations repeatable
<input checked="" type="checkbox"/> Separate for each env (dev.tfvars, prod.tfvars)	Easier testing, fewer mistakes
<input checked="" type="checkbox"/> Don't put secrets here	Use environment variables or Vault
<input checked="" type="checkbox"/> Use .auto.tfvars for modular layering	Clean and scalable
<input checked="" type="checkbox"/> Avoid mixing too many methods	Reduces confusion

---

## Summary Table

Method	Auto-loaded?	Suitable For	Precedence
-var	✗	Overrides	1
-var-file	✗	CI/CD, environments	2
Env vars	✗	Secrets, CI	3
terraform.tfvars	✓	Default settings	4
*.auto.tfvars	✓	Layered config	5
Default in code	✓	Fallbacks	6

## ✓ What Are Terraform Output Values?

**Output values** in Terraform are used to **export information** from your configuration after a successful apply, such as:

- Resource attributes (e.g., public IP, DNS name, VPC ID)
- Module results (outputs from one module to another)
- Values needed for automation or CI/CD pipelines

They make your infrastructure more **observable**, **testable**, and **integrated** with other tools or environments.

## 🔧 Syntax of an Output

```
output "<name>" {
    value      = <expression>
    description = "optional - describe the output"
    sensitive   = true|false
}
```

## Simple Example

```
resource "aws_instance" "web" {  
    ami      = "ami-0c55b159cbfafe1f0"  
    instance_type = "t3.micro"  
}  
  
output "web_instance_id" {  
    value      = aws_instance.web.id  
    description = "The EC2 instance ID"  
}
```

 After terraform apply, Terraform prints:

sh

Outputs:

```
web_instance_id = "i-0abc12345678defgh"
```

---

## Sensitive Outputs

Mark outputs as sensitive if they contain secrets like passwords or tokens.

```
output "db_password" {  
    value      = aws_db_instance.example.password  
    sensitive = true  
}
```

 Terraform will **not display** sensitive values in CLI output or logs.

---

## Real-World AWS Examples

---

◆ **1. Output EC2 Public IP for SSH Access**

```
output "ec2_public_ip" {  
    value      = aws_instance.app.public_ip  
    description = "Public IP of the application server"  
}  


---


```

◆ **2. Output Load Balancer DNS for Frontend URL**

```
output "app_url" {  
    value      = aws_lb.frontend.dns_name  
    description = "URL to access the frontend"  
}  


---


```

◆ **3. Output Subnet IDs from a VPC Module**

```
output "private_subnet_ids" {  
    value      = module.vpc.private_subnets  
    description = "Private subnet IDs used for backend"  
}  


---


```

◆ **4. Export RDS Endpoint and Port**

```
output "rds_endpoint" {  
    value      = aws_db_instance.main.endpoint  
    description = "RDS endpoint for application connectivity"  
}  
  
output "rds_port" {  
    value      = aws_db_instance.main.port
```

```
}
```

---

## ◆ 5. Output EKS Cluster kubeconfig

```
output "kubeconfig" {  
    value  = aws_eks_cluster.main.kubeconfig[0]  
    sensitive = true  
}
```

---

## Using Outputs Between Modules

You can **pass outputs from one module** into another using **module blocks**.

---

### Example: VPC Output Used in EC2 Module

#### In vpc/main.tf

```
output "vpc_id" {  
    value = aws_vpc.main.id  
}
```

#### In main.tf (root module)

```
module "vpc" {  
    source = "./vpc"  
}  
  
module "ec2" {  
    source = "./ec2"  
    vpc_id = module.vpc.vpc_id  
}
```

---

## Access Outputs in CLI & Automation

### View Outputs in CLI

sh

```
terraform output
```

```
terraform output ec2_public_ip
```

### Get Output in Scripting (e.g., Bash)

bash

```
ip=$(terraform output -raw ec2_public_ip)
```

```
echo "Server IP: $ip"
```

---

## Output with depends\_on

If your output depends on a resource's creation:

```
output "api_gateway_url" {  
    value      = aws_api_gateway_deployment.example.invoke_url  
    depends_on = [aws_api_gateway_stage.example]  
}
```

---

## Secure Handling in CI/CD (e.g., GitHub Actions)

### Masking sensitive values:

yaml

```
- name: Get RDS password  
  run: echo "::add-mask::$(terraform output -raw rds_password)"
```

---

## Useful Terraform CLI Flags with Outputs

Command	Description
terraform output	Shows all outputs (masked if sensitive)
terraform output <name>	Shows single output
terraform output -json	Returns all outputs in JSON (for scripting)
terraform output -raw <name>	Gets the raw value (no quotes, formatting)

---

### 🔥 Best Practices for Experienced DevOps Engineers

Practice	Reason
Use descriptive names	Helps others understand the infrastructure layout
Use outputs for cross-module communication	Makes modules composable and reusable
Always document outputs	Makes your infrastructure self-explanatory
Avoid exposing secrets	Use sensitive = true for passwords, tokens, etc.
Leverage outputs in pipelines	Bridge Terraform with CI/CD stages
Use terraform output -json	For structured scripting in automation